Data Streams: Probabilistic Counting

Mining Massive Datasets
Prof. Carlos Castillo
Topic 25

Sources

- Mining of Massive Datasets (2014) by Leskovec et al. (chapter 4)

  – Slides part 1, part 2

- Tutorial: Mining Massive Data Streams (2019) by Michael Hahsler

Probabilistic counting

Motivating example
how many neighbors?

- Let $n(u,h)$ be the number of nodes reachable through a path of length up to h from node u

- Naïve method

  - Maintain a set for each node u, initialize $S(u) = \{u\}$

  - Repeat h times:

$$S(u) = S(u) \cup \bigcup_{v \text{ neighbor of } u} S(v)$$

  - Answer $n(u,h) = |S(u)|$

4

What is the problem with this method?

- Let $n(u,h)$ be the number of nodes reachable through a path of length up to h from node u

- Naïve method

  - Maintain a set for each node u, initialize $S(u) = \{u\}$

  - Repeat h times:

  $$S(u) = S(u) \cup \bigcup_{v \text{ neighbor of } u} S(v)$$

  - Answer $n(u,h) = |S(u)|$

Let's look at each node

- We will receive a stream of items
  - Our neighbors at distance <= h
  - Repeated many times because of loops
- We want to use a small amount of memory
- We don't care which items are in the stream
- We just want to know **how many are distinct**

Counting fishes with pebbles

- Normally, to count with pebbles, you add one pebble every time you see an event

- How do you extend this method to count up to 1000 fishes with 10 pebbles?

- Assume you have access to a random number generator but not to an abacus for … reasons

Answer in
Nearpod Open-Ended Question

# Morris' probabilistic counting (1977)

x ← 0

For each of the n events:

x ← x + 1 with probability $(1/2)^x$

Return estimate n' = 2^x+1

Counter x needs only $log_2(n)$ bits

exact n

estimate
n' = $2^x − 1$

counter x

Simulation results by Flajolet (1985)

n (actual value)

n' (approximation)

0

Morris' algorithm provides an unbiased estimator

- Init x=0, let $p_x = 2^{-x}$, estimate $n' = 2^x - 1$

- $n = 1$
  - before: $x = 0$ $p_0 = 1$;
  - prob. 1: $x \to 1$
  - estimate $n' = 2^1 - 1 = 1 = n$

- $n = 2$
  - before: $x = 1$; $p_1 = \frac{1}{2}$
  - prob. $\frac{1}{2}$: x stays at 1; $n' = 2^1 - 1 = 1$
  - prob. $\frac{1}{2}$: $x \to 2$. $n' = 2^2 - 1 = 3$
  - $E[n'] = \frac{1}{2} \times 1 + \frac{1}{2} \times 3 = 2 = n$

Morris' algorithm provides an unbiased estimator (cont.)

Let $X(n)$ denote random counter x after n^{th} arrival

Initialize $X(0) = 0$; increment w.p. $p_x = 2^{-x}$

Estimate $n' = 2^{X(n)} - 1$

$$
\begin{aligned}
E[2^{X(n)}] \quad &= \Sigma_{j=1,\ldots,n-1} \Pr[X(n-1) = j] \, E[2^{X(n)} \mid X(n-1) = j] \\
&= \Sigma_{j=1,\ldots,n-1} \Pr[X(n-1) = j] \, (\, p_j \, 2^{j+1} + (1 - p_j) \, 2^j) \\
&= \Sigma_{j=1,\ldots,n-1} \Pr[X(n-1) = j] \, (2^j + 1) \\
&= E[2^{X(n-1)}] + 1
\end{aligned}
$$

Iterating: $E[2^{X(n)}] = E[2^{X(0)}] + n = 1 + n$

Therefore: $E[2^{X(n)} - 1] = n$

Flajolet-Martin algorithm
for counting distinct elements

- For every element u in the stream, compute hash $h(u)$

- Let $r(u)$ be the number of trailing zeros in hash value
  - Example: if $h(u) = $ `001011101000` then $r(u) = 3$

- Maintain $R = max\ r(u)$ seen so far

- Output 2^R as an estimator of the number of distinct elements seen so far

Flajolet-Martin algorithm
(intuitive, hand-waving explanation)

- Let $r(u)$ be the number of trailing zeros in hash value, keep $R = max\ r(u)$, output 2^R as estimate

- Repeated items don't change our estimates because their hashes are equal

- About ½ of distinct items hash to *******0
  - To actually see a *******0, we expect to wait until seeing 2 distinct items

- About ¼ of distinct items hash to ******00
  - To actually see a ******00, we expect to wait until seeing 4 items

- ...

- If we actually saw a hash value of ***000...0 (having R trailing zeros) then on expectation we saw 2^R different items

Flajolet-Martin, correctness proof

- Let m be the number of distinct elements
- Let $z(r)$ be the probability of finding a tail of r zeroes
- We will prove that

  $z(r) \to 1$ if $m \gg 2^r$

  $z(r) \to 0$ if $m \ll 2^r$

- Hence 2^r should be around m

Flajolet-Martin, correctness proof (cont.)

- Probability a hash value ends in r zeroes $= (1/2)^r$
  - Assuming $h(u)$ produces values at random
  - Prob. random binary ends in r zeroes $= (1/2)^r$

- Probability of seeing m distinct elements and NOT seeing a tail of r zeroes $= (1 - (\frac{1}{2})^r)^m$

Flajolet-Martin, correctness proof (cont.)

- Probability of seeing m distinct elements and NOT seeing a tail of r zeroes = $(1 - (½)^r)^m$

- Remember $(1-\varepsilon)^{1/\varepsilon} \simeq 1/e$ for small ε

- Hence

$$\left(1 - \left(\frac{1}{2}\right)^r\right)^m = \left(1 - \left(\frac{1}{2}\right)^r\right)^{\frac{m\left(\frac{1}{2}\right)^r}{\left(\frac{1}{2}\right)^r}} \approx \left(\frac{1}{e}\right)^{\left(\frac{m}{2^r}\right)}$$

Flajolet-Martin, correctness proof (cont.)

- Probability of seeing m distinct elements and NOT seeing a tail of r zeroes $\approx (1/e)^{\left(\frac{m}{2^r}\right)}$

- If $m \gg 2^r$, this tends to 0
  - We almost certainly will see a tail of r zeroes

- If $m \ll 2^r$, this tends to 1
  - We almost certainly will not see a tail of r zeroes

- Hence, 2^r should be around m

Flajolet-Martin: increasing precision

- Idea: repeat many times or compute in parallel for multiple hash functions

- How to combine?
  - Average? $E[2^r]$ is infinite, extreme values will skew the number excessively
  - Median? 2^r is always a power of 2

- Solution: group hash functions, take median of values obtained in each group, then average across groups

Let's go back to counting neighbors

Naïve method:
Maintain a set for each node u, initialize $S(u) = \{u\}$
Repeat h times: $S(u) = S(u) \cup \bigcup_{v \text{ neighbor of } u} S(v)$

Answer $n(u,h) = |S(u)|$

ANF method:

$$// \text{ Set } \mathcal{M}(x,0) = \{x\}$$
FOR each node x DO
$\quad M(x,0) = \quad$ concatenation of k bitmasks
$\qquad\qquad\qquad\qquad$ each with 1 bit set $(P(\text{bit } i) = .5^{i+1})$
FOR each distance h starting with 1 DO
$\quad$ FOR each node x DO $M(x,h) = M(x,h-1)$
$\quad // \text{ Update } \mathcal{M}(x,h) \text{ by adding one step}$
$\quad$ FOR each edge (x,y) DO
$\qquad M(x,h) = (M(x,h) \text{ BITWISE-OR } M(y,h-1))$
$\quad // \text{ Compute the estimates for this } h$
$\quad$ FOR each node x DO
$\qquad$ Individual estimate $I\hat{N}(x,h) = (2^b)/.77351$
$\qquad$ where b is the average position of the least zero bits
$\qquad$ in the k bitmasks

Palmer, C. R., Gibbons, P. B., & Faloutsos, C. (2002, July). ANF: A fast
and scalable tool for data mining in massive graphs. In Proc. KDD.

Example of another variant of the same type of algorithm

- More repetitions of the algorithm yield better precision

Becchetti, Luca, Carlos Castillo, Debora Donato, Stefano Leonardi, and Ricardo Baeza-Yates. "Using rank propagation and probabilistic counting for link-based spam detection." In Proc. of WebKDD, 2006.

21

Summary

Things to remember

- Probabilistic counting algorithms:
  - Morris
  - Flajolet-Martin

Exercises for TT22-T26

- Mining of Massive Datasets (2014) by Leskovec et al.
  - Exercises 4.2.5
  - Exercises 4.3.4
  - Exercises 4.4.5
  - Exercises 4.5.6